Model routing cuts task-call cost without detectable semantic quality loss

Same 32 Open WebUI background tasks; blinded judge scored task compliance, relevance and usefulness on a 1–5 scale.

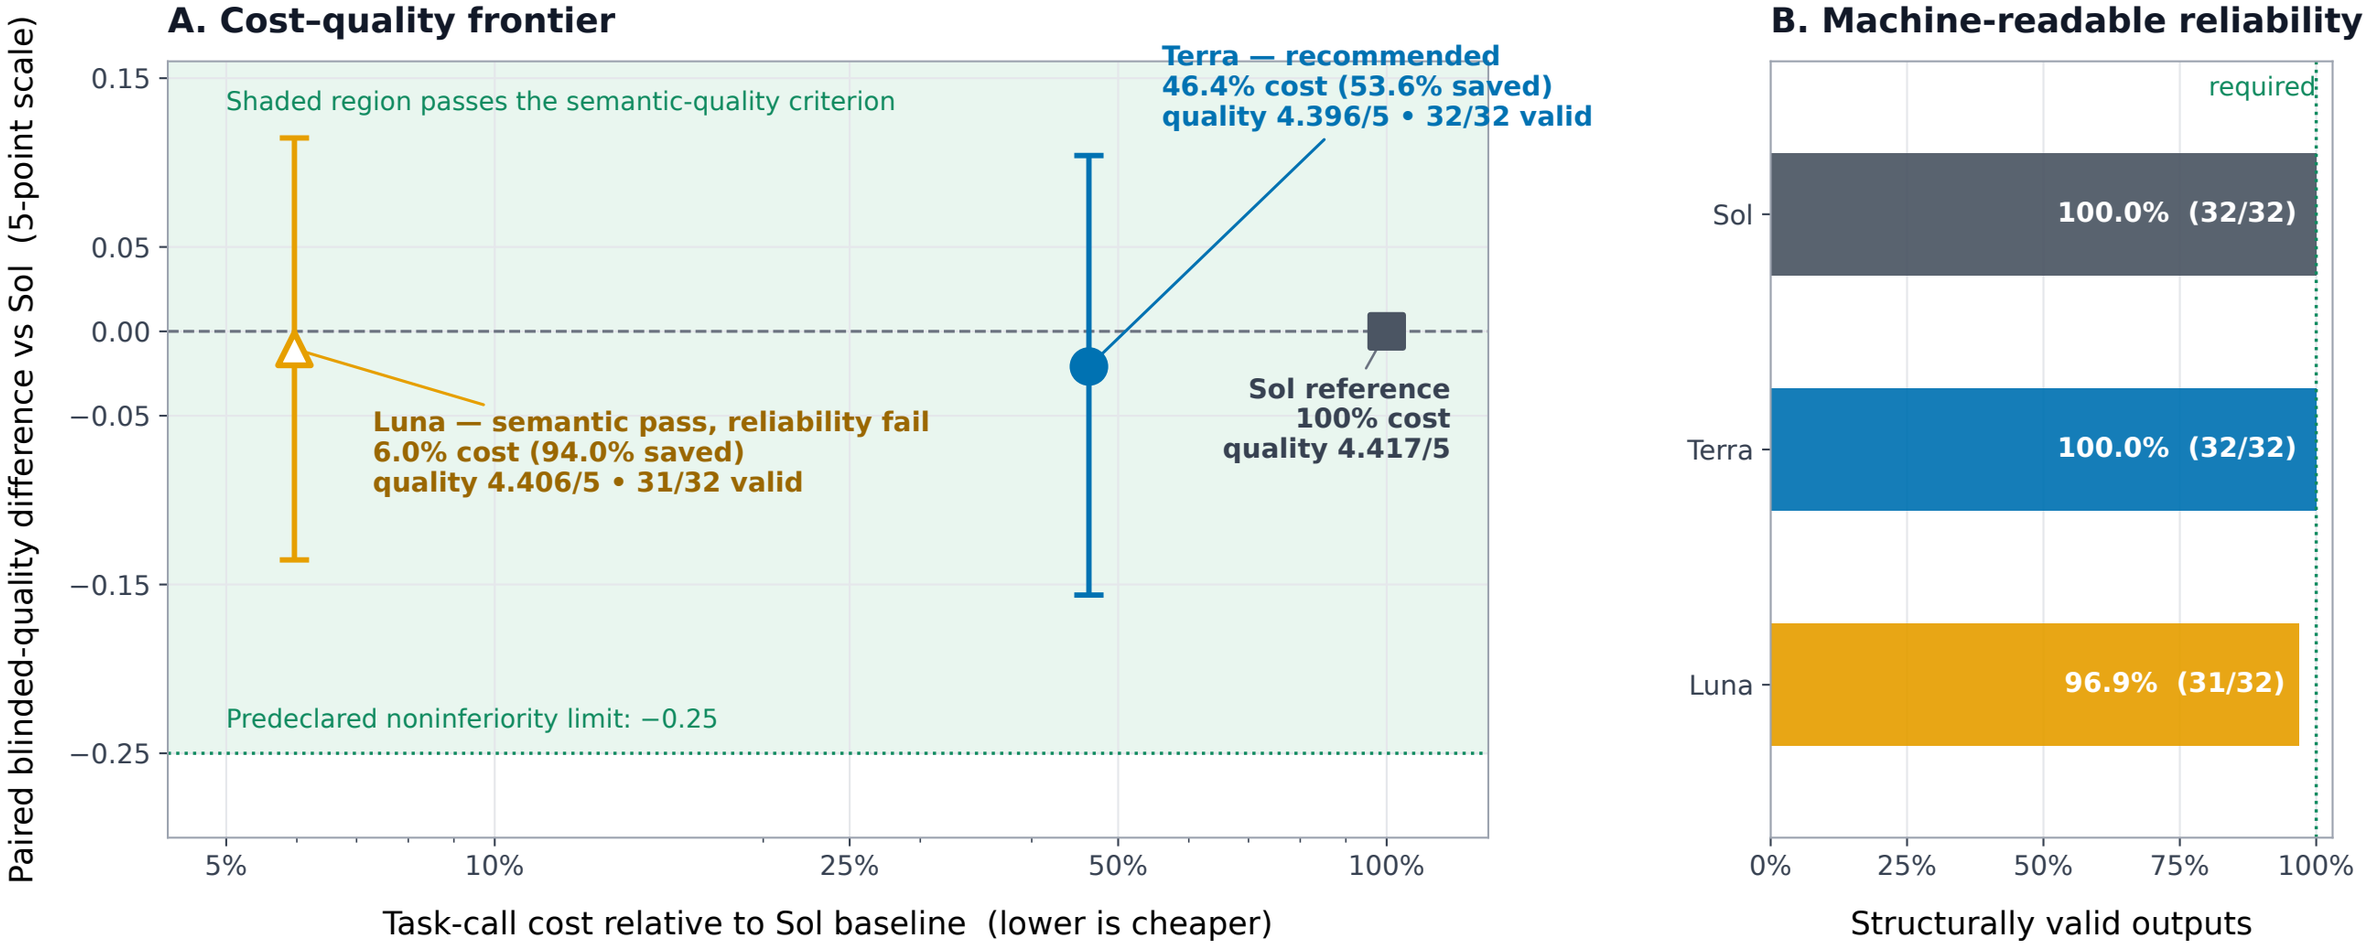

Notes: Error bars are 95% paired bootstrap intervals (10,000 resamples). Terra cost in panel A is the judged run; its three-run mean saving was 53.95% (SD 0.29 pp). Luna failed the 100% s

Cost uses provider-reported tokens and the frozen 2026-09-14 price snapshot. Main-chat responses were not rerouted. Judge: gpt-4o-2024-11-20, blinded A/B/C ordering.